
Mosaic

RUNNING META-BLOCKCHAINS TO SCALE DECENTRALIZED APPLICATIONS

**B. Bollen, P. Gevorgyan, S. Khedar,
D. Kumar Nath, M. Schenck**
for OpenST Foundation
incomplete working draft

Abstract

This is a working draft summary of notes and whiteboard sessions as we have been iterating towards the implementation for Mosaic v0.10. This is not the version intended for publication, and you are invited to review critically all content. The code is work in progress at github.com/openstfoundation/mosaic-contracts.

1. Introduction

Mosaic is a parallelisation schema for decentralized applications. Mosaic introduces a set of Byzantine fault-tolerant (BFT) consensus rules to compose heterogeneous blockchain systems into one another. Decentralized applications can use Mosaic to compute over a composed network of multiple blockchain systems in parallel.

A decentralized application is an application for which the computation is requested, performed, and paid for by independent actors. To ensure correct execution of decentralized computations first an order on the inputs of execution must be determinable. To date Ethereum is the leading network of nodes that enables developers to write and deploy code which can be called by independent users to collectively construct a shared state of the application.

The Ethereum network achieves agreement on the collective state by constructing a blockchain. A blockchain derives correctness from full replication of the computation by nodes and for Ethereum today the chain is appended through Proof-of-Work consensus. Proof-of-Work (PoW) introduces a block reward incentive for nodes to keep producing blocks and thereby securing the chain of historical blocks. However, PoW produces a serial execution model and nodes cannot be collectively rewarded for the computation they all have to perform, which renders the system computationally inefficient[1].

Active research and implementation work is ongoing to scale Ethereum's transaction throughput by dividing the verification work into multiple sections, also referred to as *shards*. At the same time Ethereum is committed to moving from the probabilistic Proof-of-Work consensus engine to a BFT based Proof-of-Stake (PoS) consensus engine.¹ The outset of a BFT consensus engine is to collectively produce a blockchain which provably – under certain honesty assumptions – cannot finalise conflicting blocks.

Thus far a vision for a decentralized web has been a major driver of innovation. However, to power global networks of billions of users and decentralized computation on vast decentralized data stores, it is reasonable to assume no single blockchain network will suffice as a backbone for all types of applications. Much like the internet is a composed network of networks, we can conceive decentralized applications to transcend any single blockchain and execute across a network of blockchain networks.

In this work we detail two protocols. The first, a layer-2 protocol, constructs *meta-blockchains* on top of existing blockchain networks to extend their state space and transaction throughput capacity. The second

¹https://github.com/ethereum/eth2.0-specs/blob/master/specs/casper_sharding_v2.1.md

protocol, called a *gateway*, acts as a message bus to send typed messages atomically between the underlying layer-1 blockchain and the meta-blockchain running on top of it.

Together these two building blocks can be used by decentralized applications to construct a parallelisation schema to increase computational capacity at lower transaction cost. In its simplest form, a decentralized application can offload transaction processing to a single meta-blockchain. More advanced applications can deploy logic across many meta-blockchains and process transactions in parallel while they can send asynchronous messages across meta-blockchains.

2. Composing Auxiliary Systems Into An Origin

A meta-blockchain composes transactions executed on an auxiliary blockchain system into an origin blockchain, such that the capacity of the origin blockchain extends and the auxiliary blockchain inherits the security properties of the origin blockchain.

To this end consider two blockchain systems, an origin blockchain O with state transition rules t_O which progress the state Σ_O

$$t_O : \Sigma_O \rightarrow \Sigma'_O \quad (1)$$

and similarly consider an auxiliary system A with state transition rules t_A such that $t_A : \Sigma_A \rightarrow \Sigma'_A$.

Transition rules for the origin and auxiliary system can be similar but do not have to be. In our discussion, when we need an example, we will reference to the origin blockchain as Ethereum (Byzantium), running Proof-of-Work, and for the auxiliary system we consider Go-Ethereum, running “clique” Proof-of-Authority (PoA). We deliberately use PoA for our considerations for the auxiliary system to emphasise that the security for the composed auxiliary system is not derived from the consensus engine of the auxiliary system. Rather the security properties of the composed system must rely only on the security properties of the consensus rules of the origin blockchain and on the Mosaic BFT consensus rules composing the auxiliary system into the origin system.²

On the origin blockchain we define a meta-blockchain \mathcal{B} with a staked validator set $\mathcal{V}_{\mathcal{B}}$. The blocks B_i of \mathcal{B} are committed to a *core* contract on origin O . For a given history of O ,³ the meta-blockchain cannot fork if we enforce that block B_{n+1} can only be (proposed and) committed after block B_n has been committed.

We define a meta-block $B_i(K, T, S) : \Sigma_A \rightarrow \Sigma'_A$ to progress the state of the auxiliary system A , where we call K the *kernel*, T the *transition object* and S the *seal*. A block B_i at height i of the meta-blockchain is committed with a seal S when a $+2/3$ supermajority of the weighted votes has been verified.⁴ We will detail later how the meta-block is constructed, but first we will describe the vote messages which, combined, seal a meta-block.

2.1. Checkpoint Finalisation - Casper FFG

The reader is assumed at this point to be familiar with the work presented in *Casper the Friendly Finality Gadget*[2] (FFG), as we will build on the logic and proofs presented there. We will intend to align with definitions and notations where possible and highlight deviations where appropriate. This section briefly (and incompletely) summarises core concepts from the paper, so that the shared concepts are established for the reader.

Casper FFG presents an algorithm to build up an overlay network of vote messages presented to a smart contract about the blocks already produced by the underlying network of nodes running the blockchain. The overlay network aims to repeatedly economically finalise a single branch of the underlying network. It does so by allowing staked validators to cast votes which together can construct a *justified chain*.

²The block proposers of the auxiliary system do have the ability to censor transactions from the auxiliary system if they have an ability to collude.

³The origin blockchain O might be probabilistically finalised, in which case a transaction to a contract can always assert that it is valid only on the intended branch of history by referencing a historical blockhash.

⁴While a vote counted on a BFT system itself is already BFT for a simple majority, we still require a supermajority, because the votes to commit a meta-block will be collected on the auxiliary system for which there is no assumption it is BFT, and we want Mosaic to be able to finalise transactions on the auxiliary system before they are committed to origin.

A justified chain is formed by a sequence of *supermajority links*, where a link identifies a *source* blockhash s and its height h_s and a *target* blockhash t and its height h_t . A vote message in Casper FFG is a link $\langle s, t, h_s, h_t \rangle_v$ signed by a validator v . A $+2/3$ supermajority signing of the same link makes it a supermajority link which justifies the target block if the source block is itself already justified.⁵

In the work it is shown that checkpoints along a justified chain can additionally be considered *finalised*, if and only if they are justified themselves and their direct child checkpoint is justified. It is then shown that, should validators finalise a checkpoint on a contradicting branch of the history of the underlying blockchain, minimally $1/3$ of the validator weight must have signed vote messages that violate either one of two rules: *the Casper slashing conditions*.

These slashing conditions can be intrinsically validated given two signed vote messages from a validator v . As such a validator must never sign two vote messages $\langle s_1, t_1, h_{s_1}, h_{t_1} \rangle_v$ and $\langle s_2, t_2, h_{s_2}, h_{t_2} \rangle_v$ for which either $h_{t_1} = h_{t_2}$ or $h_{s_1} < h_{s_2} < h_{t_2} < h_{t_1}$ holds.

It is additionally shown that validators can always finalise a new checkpoint, without being required to violate the slashing conditions, given that the block proposers propose blocks to append to the finalised fork, i.e. follow the fork selection rule of the overlay network.

2.2. Finalising Auxiliary Systems

Given a validator set \mathcal{V}_B staked for the meta-blockchain on origin, validators v can submit, on the auxiliary system, vote messages about the auxiliary system of the form

$$\langle \tilde{T}, s, t, h_s, h_t \rangle_v \quad (2)$$

where \tilde{T} is the keccak256 hash of the transition object T .

We introduce a transition object hash in the vote message to externalise otherwise intrinsic properties of the auxiliary system. This transition object allows the Mosaic validators to coarse-grain and abstract the auxiliary system into a meta-blockchain when proposing meta-blocks on origin.

We require that any transition object is calculable by a smart contract on the blockchain in question for any finalised checkpoint along a justified chain. For a given link, we define the transition object to refer to the state of the blockchain at source block s . It then follows that for a given source blockhash s , \tilde{T} is uniquely determined by s and the same properties of accountable safety and plausible liveness hold for a finalised checkpoint on a justified chain constructed by such *externalised vote messages* (2), if we extend the slashing conditions to accommodate for the transition object hash.

A validator $v \in \mathcal{V}_B$ must never sign two externalised vote messages $\langle \tilde{T}_1, s_1, t_1, h_{s_1}, h_{t_1} \rangle_v$ and $\langle \tilde{T}_2, s_2, t_2, h_{s_2}, h_{t_2} \rangle_v$ such that any of the following three conditions holds:

1. $h_{t_1} = h_{t_2}$,
2. $h_{s_1} < h_{s_2} < h_{t_2} < h_{t_1}$,
3. $s_1 = s_2 \wedge T_1 \neq T_2$.

As the slashing conditions can be intrinsically asserted to have been violated given two externalised vote messages by the same validator, there is no communication overhead to assert a possible violation of these conditions on the origin blockchain; even if the justified chain is being constructed on the auxiliary system. As a result the validators in \mathcal{V}_B can be held accountable on the origin blockchain for their voting actions on the auxiliary system. The slashing condition can be asserted on both systems and there is a clear incentive for the honest validators to assert any violation without delay on both the origin and auxiliary system, naturally synchronising the validator weights when such an event occurs.

We will later in this work address a dynamic validator set \mathcal{V}_B , where validators can join and log out on the origin system. However, already with a static validator set, we can observe that the finality gadget first introduced for (re)defining economic finality in layer-2 of Ethereum's PoW – turning miners into mere block proposers and introducing a PoS (partial) consensus engine in the smart contract layer – can also be applied to finalise an independent (auxiliary) system with a validator set whose stake is held on an external (origin) blockchain.

⁵The genesis block of the auxiliary blockchain is considered justified by definition.

2.3. Observing Origin

Per construction the finalisation of the auxiliary system by the Mosaic validators economically binds the block proposers of the auxiliary system to the Mosaic fork selection. By finalising the auxiliary system the Mosaic validators reach consensus about the auxiliary system itself. However, to construct a meta-blockchain we will additionally require the Mosaic validators to reach consensus about their observation of the origin blockchain on the auxiliary system. This is required so that message can pass bidirectionally between the chains.

To achieve it, the Mosaic validators construct a justified chain and finalise checkpoints along it for reported blocks of the origin blockchain on the auxiliary system. The incentive structure is now reversed and the origin blockchain is in no way incentivised (nor should it be) to follow the fork selection rule of how a meta-blockchain's validator set finalised its observation of the origin blockchain.

In case the origin blockchain is probabilistically finalised, then it is always possible that the Mosaic validators of a given meta-blockchain running on top of the origin blockchain would economically finalise an observation of origin on the auxiliary system which is (later) reverted by the origin blockchain - even if they sufficiently trailed the head of origin. Note that the validators cannot revert their finalised observation, because they would be required to sign vote messages which would violate the slashing conditions.

Under this scenario we must force the meta-blockchain to halt at the highest finalised checkpoint of the auxiliary system which was still consistent with its observation of the (now reverted) history of the origin blockchain. This property can be enforced by including in the transition object of a justified chain of the auxiliary system T^A information about the finalisation of the observation of the origin blockchain T^O ,

$$T_i^A = (O_f^j, \dots),$$

where $O_f^j = f(T^O)$ is the function that returns the highest finalised block number j and blockhash of the origin system as observed by the Mosaic validators for this meta-blockchain.

The origin blockchain can inspect T^A to assert that the highest finalised observation of itself on the auxiliary system is within its current history. Should this not be the case, then the origin blockchain must reject meta-blocks containing contradictory observations.

However, origin cannot directly assert that prior to the last finalised observation, there was no checkpoint finalised from a contradictory branch of its history, as the nodes on origin should not fully verify the transactions executed on the auxiliary system. This is resolved by introducing an option to challenge a proposed finalised observation.

Assume an observed checkpoint a was finalised on a contradictory branch of origin relative to the last finalised observed checkpoint b included in T^A . Any honest node can challenge the finalisation of b on origin by presenting the finalisation of a and demonstrating that $a \notin \text{history}(b)$. Note that if validators would want to alter the finalised observation from $o \rightarrow a \rightarrow b$ to $o \rightarrow b$ they would have to produce vote messages violating the slashing conditions given the already existing vote messages.

Upon success the challenger is rewarded from the stake of the offending validators and the core contract must declare the meta-blockchain halted.

2.4. Calculating the Transition Objects

We construct a transition object T^A to coarse-grain a vast amount of transactions processed on the auxiliary system under the state transition rules t_A into a single, abstracted state transition that can be validated by the core contract on the origin blockchain (under the state transition rules of the origin blockchain t_O , eq. 1).

On the auxiliary chain with every block running parameters are tracked and these consist of the latest finalised observation of the origin blockchain $O_f^j = f(T^O)$, the *accumulated transaction root* r_i , the *accumulated gas* g_i consumed, and the current *dynasty number* d on auxiliary. The transition object T^A , however, is updated at every justified checkpoint s with height h_s and we write for a given meta-block height n :

$$T_{n,d}^A(s) = (d, O_f^{j_d}, r_d, g_d, \tilde{K}_n)$$

where \tilde{K}_n is the keccak256 hash of the kernel K_n .

A smart contract calculates the parameters that go into T^A for every block. Therefore, the validators have to report the block header of every block to the smart contract. If the reported block is within the most recent

256 blocks of the auxiliary blockchain,⁶ then the smart contract can verify its correctness by accessing the corresponding block hash. If the validators fall behind more than 256 blocks in reporting, they can report more than one block per new block in order to catch up. The smart contract will record all reports, but only mark them as valid if they build a chain that reconnects to a block hash within the most recent 256 blocks of the current branch.

Tracking of T^A begins at the genesis checkpoint. For the genesis checkpoint, the accumulated transaction root is defined as the transaction root of the block, the accumulated gas consumed equals the gas consumed in the block, and the current dynasty number is 0. For all subsequent blocks, the accumulated transaction root r_i at block height i is $\text{keccak256}(r_{i-1}, R_i)$, where R_i is the transaction root of the block at height i . The accumulated gas consumed g_i at block height i is $g_{i-1} + G_i$, where G_i is the gas consumed in the block at height i . The dynasty number equals the number of finalised checkpoints in the chain from the root checkpoint to the parent block, carrying over the definition of dynasty number as in Casper FFG [2].

Further, we introduce a constant *core identifier*, a 256-bit string where the first 12 bytes are a constant specifying the origin blockchain and the rightmost 20 bytes are determined by the smart contract address of the core contract on the origin blockchain. Rather than storing the core identifier in the transaction object, it can be included as a constant in the signing string for vote messages to ensure vote messages are valid only about the intended meta-blockchain.

2.5. Proposing Meta-Blocks On Origin

For a meta-blockchain \mathcal{B} the chain can be appended by proposing and committing new meta-blocks $B_n(K_n, T_n^A, S_n)$ in the core contract on the origin blockchain. The genesis meta-block B_0 is considered committed per definition.

The kernel K_n is a tuple $(n, \tilde{p}_n, \Delta\mathcal{V}_{\mathcal{B}_n}, gp_n, gt_n)$ fully determined by and stored in the core contract on the origin blockchain. On a given branch of the origin blockchain a meta-blockchain cannot fork. The act of committing a meta-block $B_{n-1}(K_{n-1}, T_{n-1}^A, S_{n-1})$ activates the new kernel K_n at height n and *opens* the core contract to accept proposals of the form $B_n(K_n, \cdot, \cdot)$. The kernel further specifies the *parent hash* \tilde{p}_n , the *updated validator weights* $\Delta\mathcal{V}_{\mathcal{B}_n}$ to declare new validators joining and existing validators logging out. The kernel includes a voted-upon *gas price* gp_n for the gas that will be consumed in the upcoming meta-block B_n as rewarded to the Mosaic validators.⁷

gas target Lastly gt_n is the *gas target* which introduces a forcing function on the total amount of gas that can be consumed in the meta-block B_n . The earliest finalised checkpoint at which the gas target has been consumed by the meta-block must be committed to the origin blockchain by the validators. We highlight that this is not a maximum gas limit, because it cannot be guaranteed that a checkpoint can be finalised before a hard gas limit would be surpassed. Rather validators would gradually lose part of their stake as a function of the number of finalised checkpoints they failed to commit after the gas target has been surpassed.

To show that validators can be held accountable, assume the finalised checkpoint s_d has been committed which surpassed the gas target and has dynasty number d . Now assume there exists a finalised checkpoint $s_{d'}$ which also has surpassed the gas target but has a lower dynasty number $d' < d$, but was not committed before s_d . Any honest node can present the $s_{d'}$ and the core contract can assert the gas target had been consumed at a lower dynasty number than what is committed and reward the challenger.

committing a meta-block Given a committed meta-block B_{n-1} at height $n-1$ a proposal for a meta-block $B_n(K_n, T_n^A, \cdot)$ at height n is a valid proposal if validity assertions for the transition object T_n^A hold. These validity assertions require that the dynasty number and the gas consumed are strictly increasing compared to the transition object committed T_{n-1}^A . The transition object T_n^A must reference the correct kernel \tilde{K}_n . It must be checked that the latest finalised observation of origin is within the history of the current state. Furthermore

⁶This is specific to the Ethereum virtual machine, but the logic can be applied for other blockchain systems as well.

⁷A meta-blockchain has a double gas market. First the known gas market exists where users set a gas price in the transaction that gets executed on the auxiliary blockchain and the gas rewards go to the block proposers of the auxiliary system. A second, new gas market is created by requiring the block proposers to pre-deposit gas rewards for several meta-blocks in advance. Mosaic validators are rewarded upon committing a meta-block on the origin blockchain for the total gas consumed in that meta-blockchain at the agreed-upon gas price gp_n for that meta-block B_n . Details in later section.

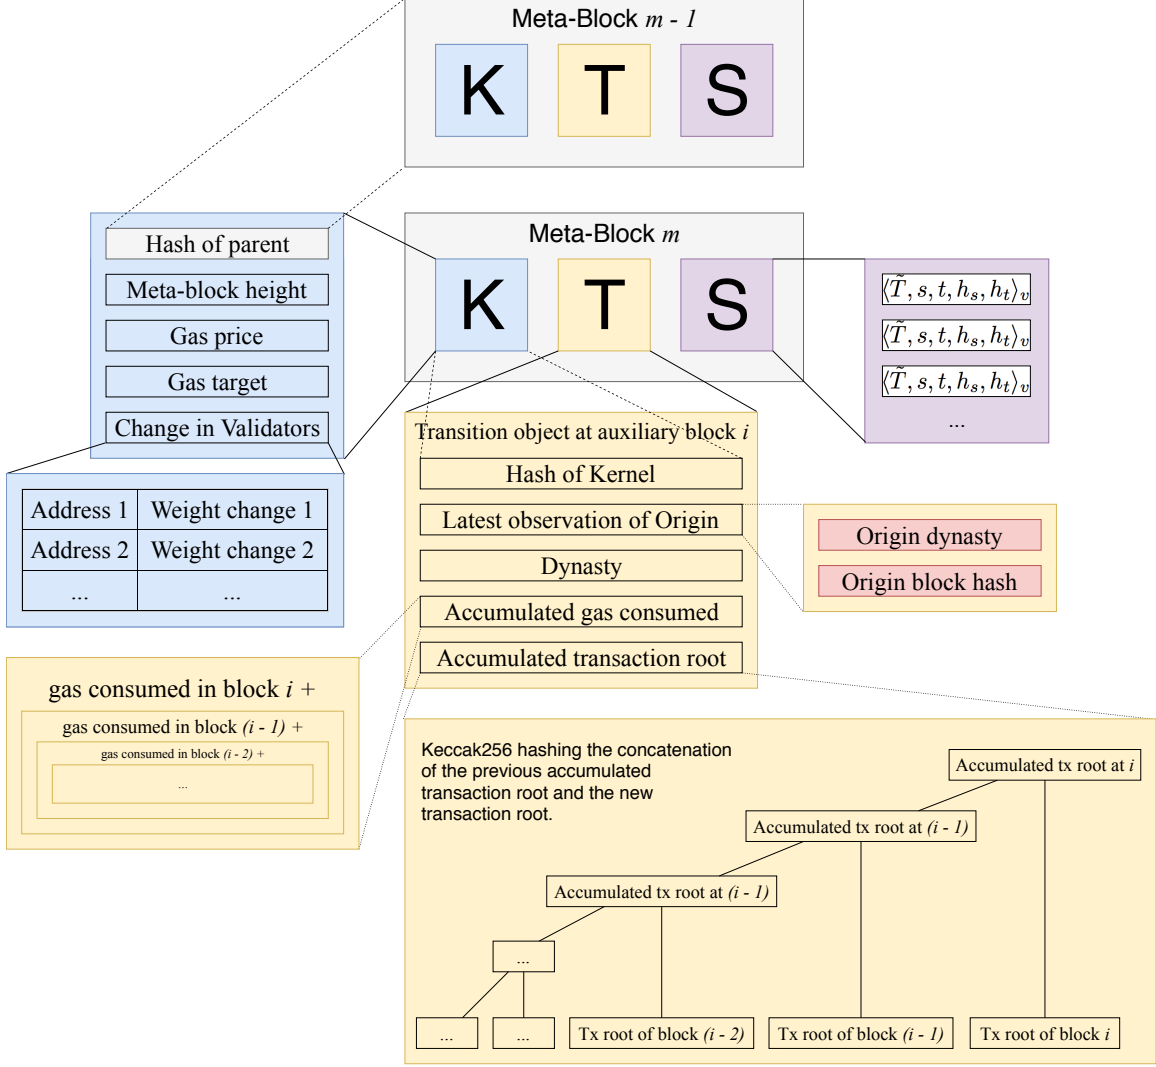


Figure 1: **Anatomy of a meta-block.** The meta-block consists of a kernel, a transition object, and a seal.

a brief challenge period exists where the transition object can be contested on the grounds that it contains a prior observation of origin that contradicts the current history of origin, as explained above in section (2.3).

It now follows that for a given active kernel K_n multiple proposals $T_{n,d}^A$ can be submitted to the core contract for different dynasty numbers d . These proposals are no longer contradicting versions of the history, rather they are ordered, sequenced snapshots of possible state transitions $B_{n,d_n}(\Sigma_{A,d_{n-1}})$ that can be committed onto origin as a new meta-block.

To commit a meta-block then requires selecting among the valid proposals a dynasty number d_n for which the meta-block will be closed. This selection occurs automatically by sequentially verifying the vote messages on the origin blockchain. The meta-block $B_n(K_n, T_n^A, S_n)$ will be committed with the seal S_n which first asserts a $+2/3$ supermajority of the voting weight \mathcal{V}_{B_n} for a valid proposal in the set of valid proposals

$$\{S_{n,d}\} = \left\{ \left\{ \langle \tilde{T}_{n,d}^A, s, t, h_s, h_t \rangle_v \right\} \right\}$$

where it is required that the supermajority link finalises the source s . As such we require that $h_t - h_s = 1$ and the source s must have been justified.

Following the same reasoning as presented in section (2.3) we can assert whether or not the finalised checkpoint s was obtained through an unbroken justified chain leading up to the finalised checkpoint, by allowing for a challenge⁸ to be raised against a proposed transition object T^A included in a vote message $\langle \tilde{T}^A, s, t, h_s, h_t \rangle_v$.

However, recall that in the case of verifying observations of origin (2.3), origin itself can act as an arbiter because origin can rely on its own history. Our concern now is to assert that the state transition to s from the last committed state B_{n-1} is a valid state transition according to the state transition rules t_A of the auxiliary system.

proof Assume first a contradictory finalisation s' of the auxiliary system relative to s exists. This implies that more than $1/3$ of the validator weight can be slashed. The core contract on origin must declare the meta-blockchain halted when it can slash more than the safety threshold of \mathcal{V}_{B_i} at a single meta-blockchain height i .⁹

Alternatively no contradictory finalisation s' of the auxiliary system relative to s exists. Under the security model assuming a $+2/3$ supermajority of the validator weight is honest, it follows that the proposed state transition to s has been obtained through honest verification of all transactions under the state transition rules t_A .

Before continuing we summarise that the proposal mechanism to commit a meta-block is asynchronous, requires no leadership-selection and is proven to be Byzantine fault-tolerant under the security model of an honest supermajority.

2.6. Beyond An Honest Supermajority

While the origin blockchain protects assets from being double-spent, users in a meta-blockchain could see their funds stolen if more than a supermajority of the open validator set would collude to violate the transition rules t_A .

We identify at least three active research topics which can allow to push the security model radically further in an efficient way; namely that (up to) a *single* honest observer can hold all validators accountable when a proposal for committing a meta-block to the origin blockchain would violate the state transition rules t_A . It is outside the scope of the current work to explore in-depth, but they are worth summarising below as future work paths.

observer network In *A Guide to 99% Fault Tolerant Consensus*[3] V. Buterin recaptures an existing algorithm by L. Lamport (1982) to describe how a *latency-dependent* network of observers can be retrofitted on top of a *threshold-dependent* consensus algorithm - in casu our here work as derived from Casper FFG - to introduce a *strong-finality* upon the finalised history constructed by the Mosaic validators arbitrarily pushing up the fault-tolerance at the cost of requiring more observers and a longer time to eventually commit (a strongly finalised proposal).

truebit A second worthwhile path is to apply the *TrueBit*[4] protocol to challenge a state transition T^A proposed by a subset of the validators \mathcal{V}_{B_i} . The validators $\{v\}$ who have signed vote messages $\langle \tilde{T}^A, s, t, h_s, h_t \rangle_v$ can be seen as the (collective) *solvers* of the off-(of-origin-)chain computation to transition the state B_{n-1} to the proposed solution s given the set of transactions specified by the accumulated transaction root r_{h_s} . The computation at hand is the execution of the virtual machine that implements the state transition rules t_A .

If less than $1/3$ of the validator weight loses the verification game to a challenger, their stake can be taken as the reward for the challenger, on top of the shared interest of all participants in the meta-blockchain for only correct proposals to be committed to the origin chain - removing the need for a jackpot and a *forced-error*. The verification game can be seen as an extension to the slashing conditions strengthening the accountable safety of the meta-blockchain to hold a minority of the validators accountable for violating the state transition rules t_A . This is made possible because the validator stake has been externalised on the origin blockchain with an independent (assumed correct) consensus mechanism.

⁸A challenger must always put forward a sufficient stake to underwrite the cost implied by her challenge.

⁹When a meta-blockchain has been halted, assets and stateful objects can be recovered on the origin blockchain at any later time with Merkle proofs against the last committed state root of the meta-blockchain. We will detail the conditions under which a meta-blockchain must halt and how the recovery processes can work later in this writing.

If more than $2/3$ of the validator weight signed the vote message which was challenged and lost the verification game, their stake is the reward for the challenger, and the meta-blockchain must be declared halted by the core contract.

argument of knowledge A third strategy can be constructed by requiring any of the challenged validators of a proposal T^A to present an *argument of knowledge*, eg. using zkSNARKs generated after being challenged, to cryptographically prove that the block s has been obtained by correctly applying the state transition rules t_A recursively on each parent block building up from the committed block found in B_{n-1} .

The costs both in time and expense to generate this proof by a validator scale at first approximation linearly with the gas consumed in the proposed meta-block B_n . The transition object T^A to commit this meta-block tracks the gas consumed, so the core contract can require that the challenger puts forward a sufficient stake to substantiate her claim.

There are open considerations regarding the (economic) feasibility of requiring all validators to be able to generate such a proof on commodity hardware. The task of generating these proofs could be designated to a specialised node, but an incentive problem arises when it is economically unattractive to run such a node considering it is unlikely that a proposal would be challenged.

We have gone beyond the assumption of an honest supermajority of validators by considering three strategies discussed in the field to radically improve the security model; up to the point where a single honest observer can cost-effectively challenge the validator set on the origin blockchain.

2.7. Dynamic Set of Validators

The set of validators for the meta-blockchain \mathcal{B} must be dynamic to allow new validators to join and for validators to safely exit from their responsibilities. It is defined for each meta-block height i and we note $\mathcal{V}_{\mathcal{B}_i}$.

A validator v is defined as an address which has a strictly-positive validator weight w_v associated with it on the origin blockchain. A validator can join the set of validators by sending a *deposit message* to the origin blockchain and putting forward a stake. A validator can initiate log out by sending a *logout message* to origin.

The weight of a validator v at meta-block height n is the product of the stake of the validator at that height $s_v(n)$ multiplied with a reputation function $R(n_0, n; v)$

$$w_v(n) = s_v(n) \cdot R(n_0, n; v) \quad (3)$$

where the reputation function R has values in the range $[0, 1]$ and n_0 is the meta-block height where the validator v sent the deposit message.

Changing the reputation function allows for different security models to mitigate collusion attempts of malicious actors within the validator set of the meta-blockchain. As examples, the reputation function can require initial work by the validator before obtaining a non-zero reputation. The reputation function can also enforce a finite (probabilistic) lifetime for all validators, honest and dishonest alike. Inducing churn in the (honest) validator set adds an important stochastic tool to prevent collusion of more than $1/3$ of the validator weight.

Changes to the validator weights $w_v(n)$ are included in the kernel K_n as $\Delta \mathcal{V}_{\mathcal{B}_n}$. When tallying votes both the origin blockchain must use the *effective weight* $W_v(n)$,

$$W_v(n) = w_v(n) \cdot \theta_v$$

where θ_v is initially 1 for all validators, but irreversibly and immediately set to 0 when a validator v is shown to have violated any slashing condition. The effective weight allows for the voting power of an offending validator to be revoked on both blockchains instantaneously without waiting for the meta-block to be committed.

The stake $s_v(n)$ of a validator can only be preserved or decrease for increasing n when the validator does not fulfill its responsibilities. The *loss function* L recalculates the stake of a validator for the new meta-blockheight

$$s_v(n) = L(s_v(n-1); v). \quad (4)$$

The loss function reduces the actual value the validator can possibly withdraw in the future. The loss function can reward challengers for presenting proofs of bad behaviour by the validator from his stake and can burn the remaining reduction.

When a validator violates any slashing condition the loss function returns zero.

Validators can see their stake reduced when they become inactive in the voting process (*inactivity leak*), or when they fail to commit a meta-block back to origin according to the gas target requirements (see above).

When a validator successfully logs out his reputation is set to zero, but his stake remains locked in the contract for a fixed amount of time to protect against long range revision attacks.¹⁰ After this time the validator can withdraw his stake (and remaining rewards).

When a validator wants to leave the set of validators, it sends a logout message to the stake contract on origin. The validator's weight is reduced to zero, therefore the validator leaves the set of validators.




Figure 2: **Transition.** After the transition object T_i is committed on origin, the opening of the new kernel K_{i+1} is confirmed on auxiliary. On auxiliary, meta-block B_{i+1} retroactively opens from the point where B_i closed. The set of validators, however, only changes at the time of confirmation. The dashed line under auxiliary shows the active set of validators.

When a validator casts a vote on auxiliary, the vote carries a specific weight. The initial weight of a new validator is equal to the stake on origin. When we say that a vote receives a $2/3$ majority, we mean that at least $2/3$ of the total voting weight on auxiliary have cast that vote. A validator's stake may be reduced to a non-zero value, e.g. by slashing. However, in any slashing case the weight will always go to zero immediately.

We defend against long range revision attacks the same way casper does. Four months after the validator left the set of validators, the validator can withdraw its stake.

Casper FFG[2] defines a stitching mechanism for a changing set of validators. It prevents a case where a significant change in validators from one dynasty to the next could result in conflicting finalised checkpoints on different forks. They define a forward validator set and a rear validator set based on the start dynasty and end dynasty of a validator. In the case of Mosaic, however, validators do not join and leave the set of validators within two dynasties, but rather two meta-blocks. When a validator sends a deposit or logout message at meta-block height i , the change is announced to auxiliary with the opening of meta-block $i + 1$. The auxiliary system will update the validator's weight with the opening of meta-block $i + 2$ and the validator will join or leave the set of validators. We call that meta-block the start meta-block or end meta-block. Following this change, the forward and rear validator sets are based on the start meta-block and end meta-block, respectively, instead of the start dynasty and the end dynasty.

In the same way, an ordered pair of checkpoints (s, t) , where t is in meta-block B_i , has a supermajority link if both at least $2/3$ of the forward validator set and the rear validator set of B_i , respectively, have published vote $\langle \tilde{T}, s, t, h_s, h_t \rangle_v$.

The same public key can only deposit and withdraw once. It cannot join the set of validators multiple times or change its stake.

NOTE: weight = stake x reputation

¹⁰By holding on to the stake the validator can still be held accountable for evidence of wrong doing that may surface later or new attempts of revising the history. We follow Casper FFG's proposal of approximating four months delay before withdrawal is allowed.

When a meta-block has been committed on origin, the Kernel of the newly opening meta-block can be confirmed on auxiliary. The state root of origin is regularly transferred to auxiliary and finalised. By means of a Merkle proof against that known state root, the content of the Kernel can be made available to auxiliary. Auxiliary will note the update of the validator weights for the meta-block height of the opened meta-block plus one. The validator set changes on auxiliary at the time the kernel opening is confirmed. Even though the kernel opening retroactively opens the meta-block on auxiliary from the first block after the last finalised checkpoint in the meta-block that precedes the opening one, the validator set only changes at the point of the confirmation. When the Kernel at height i is confirmed on auxiliary, the validator set $i + 1$ becomes the active set of validators. The validators that logged out on origin before the opening of K_i will be in the rear validator set of K_{i+1} and leave the validator set when K_{i+2} is confirmed.

Figure 3: **Dynamic set of validators.** The validator sends a log out message (or a deposit message) in B_i . The change in the set of validators is transferred to auxiliary as part of kernel K_{i+1} that gets confirmed. The validator does its last round in the rear validator set on auxiliary. The validator’s weight drops to zero when K_{i+2} is confirmed (or the validator joins in the forward set of validators).

Users of the auxiliary system pay gas fees to the block proposers of the auxiliary system for every transaction. Similarly, users of the origin blockchain pay gas fees for their transactions on origin. Due to the limit on the number of transactions included in each block, the gas price is determined in a system of supply and demand. When many users want to include many transactions in blocks, possibly above the limit, the average gas price will naturally rise and vice versa.

Mosaic provides two incentives for the validators to do so:

Gas fees. For every meta-block, validators get rewards based on the amount of gas that was consumed on auxiliary within that meta-block.

where w_v is the weight of the validator, w is the total weight of all validators, g is the accumulated gas of the meta-block, and gp is the agreed upon gas price.

However, this amount increases with increasing size of a meta-block.

Validator rewards. Whenever someone commits a meta-block on origin, the meta-block closes, a new meta-block opens, and someone confirms the kernel of the newly opened meta-block on auxiliary. The auxiliary system tracks the total gas consumption per checkpoint and can therefore calculate the amount of gas consumed in the closed meta-block on-chain. The gas consumed of a meta-block is merely a mirror of the gas that all auxiliary transactions within that meta-block consumed on auxiliary.

A smart contract on auxiliary pays out a fee to the validators based on the gas consumed in the closed meta-block. In order for the auxiliary smart contract to be able to do the pay-out, the block proposers of auxiliary must deposit sufficient value for multiple future meta-blocks in said contract. Block proposers have an interest in confirmed meta-blocks on origin. They earn on auxiliary for proposing blocks and they can only transfer their value to origin after a meta-block is committed. Furthermore, an auxiliary system with more frequent meta-block commits could be more interesting to users and therefore produce more gas rewards for the block proposers.

If block proposers would fail to produce a sufficient deposit for validator rewards, validators would most likely switch to a more profitable auxiliary chain.

3. A Mosaic of Cores

summary: by constructing asynchronous BFT consensus rules to compose a heterogeneous auxiliary blockchain into the origin blockchain, we leverage the security of the origin blockchain to asynchronously finalise the transactions on the auxiliary blockchain. As there are no time constraints or time-outs in the Mosaic consensus rules, the limited processing capacity of the origin blockchain does not undercut the ability to compose many auxiliary systems into the same origin blockchain. This enables Mosaic to run many meta-blockchains in parallel, additively improving the total transaction capacity, but keeping a bound on the cost imposed on the origin blockchain.

4. Gateway

abstract: the gateway protocol allows decentralized applications to deploy a message bus to send asynchronous messages between the state space of the origin blockchain and the state space of the auxiliary system supporting a specific meta-blockchain. This allows tokens (ERC20, later non-fungible tokens) declared on the origin blockchain (Ethereum) to be locked and moved into the meta-blockchain where transactions can be finalised faster, and at a lower transaction cost.

The gateway protocol is code-complete at github.com/openstfoundation/mosaic-contracts

5. Outlook

6. Conclusion

References

- [1] Luu, L., Teutsch, J., Kulkarni, R. & Saxena, P. Demystifying incentives in the consensus computer. *IACR Cryptology ePrint Archive* **2015**, 702 (2015).
- [2] Buterin, V. & Griffith, V. Casper the friendly finality gadget. *CoRR* **abs/1710.09437** (2017). URL <http://arxiv.org/abs/1710.09437>. 1710.09437.
- [3] Buterin, V. A guide to 99% fault tolerant consensus (2018). URL https://vitalik.ca/general/2018/08/07/99_fault_tolerant.html.
- [4] Teutsch, J. & Reitwiener, C. A scalable verification solution for blockchains (2016). URL <https://people.cs.uchicago.edu/~teutsch/papers/truebit.pdf>.